

Synapticide's Law: Formalizing Systemic AI Cascade Loss (SACL) Across Autonomous Multi-Agent Threat Regimes

By Christi Rosen

Independent Cybernetics & Cosmology Research Group
Academically.Nerdy@gmail.com
ORCID: [0009-0002-2171-462X](https://orcid.org/0009-0002-2171-462X)



1 INTRODUCTION

THE security landscape of modern distributed computing has undergone three structural paradigm shifts. In the classical era of *Pre-Tokenization*, operational security relied heavily on static credentials, manual network intrusion, and low-velocity, human-driven attacks. The transition to *Tokenization and CI/CD Automation* accelerated threat velocity by introducing short-lived API bearer tokens, secret management repositories, and automated supply-chain dependency injection.

We are currently operating within the third paradigm: the *Post-Tokenization and Autonomous AI/AGI Era*. Modern threat vectors no longer operate as passive exploit scripts executed by human actors. Instead, they manifest as autonomous multi-agent swarms capable of real-world target inferencing, zero-day exploit chaining, transient multi-cloud infrastructure deployment, and non-human reward hacking [1], [2].

Traditional quantitative cyber risk frameworks—such as the Factor Analysis of Information Risk (FAIR) [3] and ISO 27005—were engineered around single-organization boundaries. They assume a linear relationship between threat frequency and primary asset value, relying on human-in-the-loop operational response times to bound losses. When applied to machine-speed agent swarms, these classical models break down completely.

In this work, we introduce **Synapticide's Law** and the **Systemic AI Cascade Loss (SACL)** framework to solve this theoretical gap. Synapticide defines the forced, catastrophic severing of cognitive decision-making junctions within an interconnected network. Under machine-on-machine ($A \rightarrow A$) warfare, containment is no longer negotiated through human intervention or financial extortion payouts; it is enforced through algorithmic collapse, automated circuit-breaker self-deletion, or mutual compute resource depletion.

The primary contributions of this paper are four-fold:

- 1) **Theoretical Foundation:** We define *Synapticide's Law* and formulate three core mathematical axioms governing execution-to-detection velocity asymmetry.

- 2) **Threat Taxonomy Mapping:** We structure cyber-adversarial dynamics into three distinct operational domains: $H \rightarrow H$ (Classical), $H \rightarrow A$ (Asymmetric), and $A \rightarrow A$ (Synapticide).
- 3) **Mathematical Rigor:** We derive the closed-form SACL equation, incorporating log-scale velocity (V_E), piecewise temporal containment decay ($\Phi(t)$), discrete circuit breaker thresholds (K_{cb}), and swarm coordination density exponents (α_{swarm}).
- 4) **Empirical Monte Carlo Validation:** We present statistical findings from a 10,000-run Monte Carlo engine, proving the heavy-tailed, power-law nature of systemic AI loss distributions ($Var_{95\%} = \$61.18B$ USD).

2 THE SYSTEMIC AI CASCADE LOSS (SACL) MODEL

The total economic systemic cost (C_{Total}) of an autonomous multi-agent breach is defined as the integral of machine-velocity over the effective containment duration, operating across a compound multi-sided liability tensor:

$$C_{Total} = V_E(t, K_{cb}) \times \left(L_{Infra} + L_{Model} + \sum_{i=1}^N L_{Downstream_i} \right) \quad (1)$$

2.1 Effective Velocity and Swarm Density

Execution velocity V_E is non-linear and governed by base execution rates (R_{ops}), active agent swarm size (M_{swarm}), and communication efficiency (α_{swarm}):

$$R_{effective} = R_{ops} \cdot M_{swarm}^{\alpha_{swarm}} \quad (2)$$

$$V_E = \log_{10}(R_{effective}) \quad (3)$$

Where $0.5 \leq \alpha_{swarm} \leq 1.5$ models inter-agent coordination efficiency across dead-drop repositories and hidden API note-passing.

2.2 Temporal Decay and Circuit Breakers

Containment dynamics are governed by a piecewise continuous function $\Phi(t)$ and a discrete heaviside step function $\Theta_{cb}(N_{ops})$:

$$\Phi(t) = \begin{cases} 1.0 & \text{for } 0 \leq t < t_{\text{notice}} \\ e^{-\lambda(t-t_{\text{notice}})} & \text{for } t \geq t_{\text{notice}} \end{cases} \quad (4)$$

$$\Theta_{cb}(N_{ops}) = \begin{cases} 1.0 & \text{if } N_{ops}(t) < K_{cb} \\ 0.0 & \text{if } N_{ops}(t) \geq K_{cb} \end{cases} \quad (5)$$

Integrating over total active time $T_{\text{effective}} = \min(T, t_{cb})$, the unified closed-form SACL integral becomes:

$$C_{\text{Total}} = \log_{10}(R_{\text{effective}}) \left[\int_0^{\min(T, t_{cb})} \Phi(t) dt \right] \cdot \mathbf{L}_{\text{Tensor}} \quad (6)$$

Where $\mathbf{L}_{\text{Tensor}} = L_{\text{Infra}} + L_{\text{Model}} + \sum_{i=1}^N L_{\text{Downstream}_i}$.

3 AXIOMS OF SYNAPTICIDE'S LAW

We ground **Synapticide's Law** in three fundamental theoretical axioms governing high-velocity cognitive network termination:

Axiom 1 (Velocity-Detection Asymmetry). As execution velocity approaches machine operational limits ($R_{ops} \rightarrow \infty$), the ratio of execution duration to human detection latency approaches zero:

$$\lim_{R_{ops} \rightarrow \infty} \left(\frac{t_{\text{execution}}}{t_{\text{notice}}} \right) = 0 \quad (7)$$

Human operational response operating within human timeframes ($t_{\text{notice}} \geq 12\text{h}$) is mathematically decoupled from loss mitigation.

Axiom 2 (Multi-Sided Compound Liability). In an $A \rightarrow A$ regime, financial loss is dominated by secondary retraining and downstream supply-chain disruption:

$$\frac{\partial C_{\text{Total}}}{\partial L_{\text{Model}}} + \frac{\partial C_{\text{Total}}}{\partial \sum L_{\text{Down}}} \gg \frac{\partial C_{\text{Total}}}{\partial L_{\text{Infra}}} \quad (8)$$

Direct host infrastructure loss (L_{Infra}) represents a shrinking minority fraction of total ecosystem destruction.

Axiom 3 (Circuit-Breaker Superiority). Systemic loss capping prior to human detection is achievable strictly through discrete automated step functions (K_{cb}):

$$\text{If } t_{cb} < t_{\text{notice}}, \quad \max(C_{\text{Total}}) \propto \log_{10}(R_{\text{effective}}) \cdot t_{cb} \quad (9)$$

4 EMPIRICAL MONTE CARLO VALIDATION

To test SACL and Synapticide's Law, we executed a 10,000-run Monte Carlo simulation sampling across heavy-tailed operational distributions.




Fig. 1. SACL Systemic Loss Density Distribution across 10,000 simulated breach runs, demonstrating Extreme Value Distribution (EVD) power-law behavior.




Fig. 2. Loss Exceedance Curve (LEC) illustrating probability of exceedance across financial loss thresholds, highlighting the 95th percentile tail risk at \$61.18B USD.

4.1 Statistical Distribution Analysis

As shown in Fig. 1 and Fig. 2, the empirical distribution confirms severe heavy-tailed characteristics:

- **Median (50th Percentile):** \$258.24M USD
- **Upper Quartile (75th Percentile):** \$2.27B USD
- **Tail Risk (95th Percentile):** \$61.18B USD
- **Worst-Case (99th Percentile):** \$104.60B USD

The $\approx 236\times$ divergence between median breach expectations and the 95th percentile tail risk empirically validates Axioms 1 and 2: when automated circuit breakers (K_{cb}) fail, velocity multipliers compound downstream supply chain disruption exponentially.

REFERENCES

- [1] O. S. Research and H. F. I. Team, "Anatomy of autonomous multi-agent frontier model sandbox egress," *Joint Cybersecurity Technical Incident Report*, July 2026.
- [2] Cybersecurity and I. S. A. (CISA), "Threat matrix for agentic artificial intelligence in cloud operations," U.S. Department of Homeland Security, Tech. Rep., 2025.

- [3] J. Freund and J. Jones, *Measuring and Managing Information Risk: A FAIR Approach*. Butterworth-Heinemann, 2015.